

Financial Timeseries Test

August 21, 2021

1 Introduction to the Task

We would like to build a model that can predict the daily closing price of US stocks. Because a company's historical price data alone does not contain all information about what the future close prices might be, please be sure to include contextual information about the company, such as industry/sector, in your analysis.

For this task, please use the following datasets: - [stocks dataset](https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/8c4c02a9a49365719dee849a8bfeb64e/Stocks2.zip) [https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/8c4c02a9a49365719dee849a8bfeb64e/Stocks2.zip](https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/193bb5ff397532fa61b2a6417a125e90) - [stocks dataset cont](https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/193bb5ff397532fa61b2a6417a125e90) <https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/193bb5ff397532fa61b2a6417a125e90> - [etfs](https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/874ba49094e7d05cef0d82fc2fd8ef1a/ETFs.z) <https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/874ba49094e7d05cef0d82fc2fd8ef1a/ETFs.z> - [company info](https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/4cfaa8631b61675dfa033b316ad3b) <https://hr-projects-assets-prod.s3.amazonaws.com/3maoh37pfp9/4cfaa8631b61675dfa033b316ad3b>

You may limit your analysis to any date range you choose. If you choose to work with a subset companies in the given dataset, please justify how you selected which companies to include.

```
[1]: # If you'd like to install packages that aren't installed by default, list them
      ↪ here.
      # This will ensure your notebook has all the dependencies and works everywhere

import sys
!{sys.executable} -m pip install xgboost sklearn
```

Requirement already satisfied: xgboost in /opt/conda/lib/python3.7/site-packages (1.4.2)

Requirement already satisfied: sklearn in /opt/conda/lib/python3.7/site-packages (0.0)

Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.19.2)

Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.3.3)

Requirement already satisfied: scikit-learn in /opt/conda/lib/python3.7/site-packages (from sklearn) (0.21.3)

Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-packages (from scikit-learn->sklearn) (0.17.0)

```
[2]: import pandas as pd
      import seaborn as sns
```

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error
```

2 Step 1: Load, merge, and explore the data

What does the data look like? What are some general characteristics? Are there any anomalies that might impact the later modeling steps?

```
[3]: companies = pd.read_csv("https://hr-projects-assets-prod.s3.amazonaws.com/
    ↪3maoh37pfp9/4cf8a8631b61675dfa033b316ad3bbe5/companies.csv")
companies.head()
```

```
[3]:  ticker                company name                short name \
0      A      Agilent Technologies Inc.      Agilent
1     AA      Alcoa Corporation      Alcoa
2    AABA      Altaba Inc.      Altaba
3     AAC      AAC Holdings Inc.      AAC
4   AADR  AdvisorShares Dorsey Wright ADR  AdvisorShares Dorsey Wright

                                industry \
0  Medical Diagnostics & Research
1                Metals & Mining
2            Asset Management
3      Health Care Providers
4                      NaN

                                description \
0  Agilent Technologies Inc is engaged in life sc...
1  Alcoa Corp is an integrated aluminum company. ...
2  Altaba Inc is an independent, non-diversified,...
3  AAC Holdings Inc provides inpatient and outpat...
4  The investment seeks long-term capital appreci...

                                website      logo      ceo \
0      http://www.agilent.com      A.png      Michael R. McMullen
1      http://www.alcoa.com      AA.png      Roy Christopher Harvey
2      http://www.altaba.com      AABA.png      Thomas J. Mcinerney
3  http://www.americanaddictioncenters.org      NaN      Michael T. Cartwright
4      http://www.advisorshares.com      AADR.png      NaN

                                exchange      market cap      sector \
```

0	New York Stock Exchange	2.421807e+10	Healthcare
1	New York Stock Exchange	5.374967e+09	Basic Materials
2	Nasdaq Global Select	4.122368e+10	Financial Services
3	New York Stock Exchange	6.372010e+07	Healthcare
4	NYSE Arca	1.031612e+08	NaN

	tag 1	tag 2	tag 3
0	Healthcare	Diagnostics & Research	Medical Diagnostics & Research
1	Basic Materials	Aluminum	Metals & Mining
2	Financial Services	Asset Management	NaN
3	Healthcare	Medical Care	Health Care Providers
4	NaN	NaN	NaN

```
[4]: #Show all categories for "exchange" to see which one is related to NASDAQ
companies["exchange"].unique()
```

```
[4]: array(['New York Stock Exchange', 'Nasdaq Global Select', 'NYSE Arca',
        'NYSE American', 'NASDAQ Global Market', 'NASDAQ Capital Market',
        'Cboe Global Markets EDGX', 'Investors Exchange', 'OTC Pink'],
        dtype=object)
```

```
[5]: # To see which NASDAQ category has the maximum count
companies["exchange"].value_counts()
```

```
[5]: New York Stock Exchange      2467
Nasdaq Global Select             1432
NASDAQ Capital Market            793
NYSE Arca                        745
NASDAQ Global Market             621
NYSE American                   273
Cboe Global Markets EDGX         35
OTC Pink                         1
Investors Exchange               1
Name: exchange, dtype: int64
```

```
[6]: # Filter just the companies that belong to "Nasdaq Global Select"
nas_df = companies[companies["exchange"] == "Nasdaq Global Select"]
nas_df.head()
```

```
[6]:   ticker      company name      short name \
2    AABA      Altaba Inc.      Altaba
5    AAL  American Airlines Group Inc.  American Airlines
10   AAON      AAON Inc.      AAON
12   AAPL      Apple Inc.      Apple
15   AAWW  Atlas Air Worldwide Holdings  Atlas Air Worldwide

      industry \
```

```

2           Asset Management
5           Airlines
10          Building Materials
12          Computer Hardware
15  Transportation & Logistics

```

```

description \
2  Altaba Inc is an independent, non-diversified,...
5  American Airlines Group Inc operates over 6,00...
10 AAON Inc is a heating, ventilation and air con...
12 Apple Inc is designs, manufactures and markets...
15 Atlas Air Worldwide Holdings Inc is engaged in...

```

```

website      logo      ceo \
2  http://www.altaba.com  AABA.png  Thomas J. Mcinerney
5  http://www.aa.com     AAL.png   W. Douglas Parker
10 http://www.aaon.com   AAON.png  Norman H. Asbjornson
12 http://www.apple.com  AAPL.png  Timothy D. Cook
15 http://www.atlasair.com  AAWW.png  William J. Flynn

```

```

exchange      market cap      sector \
2  Nasdaq Global Select  4.122368e+10  Financial Services
5  Nasdaq Global Select  1.694019e+10  Industrials
10 Nasdaq Global Select  1.961880e+09  Basic Materials
12 Nasdaq Global Select  8.074917e+11  Technology
15 Nasdaq Global Select  1.395183e+09  Industrials

```

```

tag 1      tag 2      tag 3
2  Financial Services      Asset Management      NaN
5  Industrials            Airlines              NaN
10 Basic Materials        Building Materials    NaN
12 Technology            Consumer Electronics  Computer Hardware
15 Industrials  Airports & Air Services  Transportation & Logistics

```

```

[7]: # Read the dataset from the company Altaba
df = pd.read_csv('aaba.us.txt', delimiter = ",")

```

```

[8]: df[["Date", "Close"]]

```

```

[8]:      Date  Close
0   1996-04-12   1.38
1   1996-04-15   1.34
2   1996-04-16   1.20
3   1996-04-17   1.12
4   1996-04-18   1.22
...      ...    ...
5429 2017-11-06  71.71


```

```
5430 2017-11-07 72.22
5431 2017-11-08 71.16
5432 2017-11-09 70.19
5433 2017-11-10 70.56
```

```
[5434 rows x 2 columns]
```

```
[9]: # Plot the "Close" information
plt.figure(figsize=(8, 6))
plt.plot(df["Close"])
```

```
[9]: [<matplotlib.lines.Line2D at 0x7f826b706ed0>]
```



3 Step 2: Build a Model

Given the data explored in the previous step, predict tomorrow's closing price for each company in the NASDAQ index. If it helps simplify the analysis, you may build a model that predicts closing price(s) for a subset of companies in the NASDAQ index.

```
[10]: # This is a function that transforms a series into a

def transformSeriesToDataset(series, NumberOfElements):
    dataset = None
    outputDataset = None
    for counter in range (len(series)-NumberOfElements-1):
        sample = np.array([series[counter:counter+NumberOfElements]])
        output = np.array([series[counter+NumberOfElements]])
        if dataset is None:
            dataset = sample
        else:
            dataset = np.append(dataset,sample,axis = 0)
        if outputDataset is None:
            outputDataset = output
        else:
            outputDataset = np.append(outputDataset,output)
    return dataset, outputDataset


[11]: # Create a series variable with the "Close" field

series = df['Close'].to_numpy()
X, Y = transformSeriesToDataset(series, NumberOfElements = 10)


[12]: # Split the information in train and test

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.4)


[13]: # Creation of a lineal model and fitting
lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)


[13]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)


[14]: y_train_predict = lin_model.predict(X_train)
MSE = mean_squared_error(Y_train,y_train_predict)
print("Entrenamiento: MSE =" +str(MSE))

y_test_predict = lin_model.predict(X_test)
MSE = (mean_squared_error(Y_test, y_test_predict))
print("Pruebas: MSE =      " +str(MSE))

df_predictions = pd.DataFrame({'Real_Value':Y_test, 'Prediction':
    ↳y_test_predict, 'Diff':Y_test-y_test_predict})
df_predictions = df_predictions.reset_index(drop = True)
df_predictions.head(10)
```

Entrenamiento: MSE =1.1831113909006064
Pruebas: MSE = 1.1620926887716074

```
[14]: Real_Value Prediction      Diff
0      29.64    29.333274  0.306726
1      17.31    17.472378 -0.162378
2      16.09    16.186748 -0.096748
3       8.52     8.898262 -0.378262
4      21.41    21.117678  0.292322
5      42.99    45.463895 -2.473895
6       4.33     4.472247 -0.142247
7      15.73    15.858349 -0.128349
8      11.42    10.767273  0.652727
9      23.72    23.612139  0.107861
```

4 Step 3: Evaluate the model on test data

How does the model perform on data that it hasn't seen? When predictions are poor, why does the model fail to predict the close price?

```
[15]: # We will do the test in the American Airlines dataset, which belongs to the
      ↪ "Nasdaq Global Select" exchange
df2 = pd.read_csv('aal.us.txt', delimiter = ",")
df2[["Date", "Close"]].head(15)
```

```
[15]:      Date      Close
0  2013-12-10  24.064
1  2013-12-11  25.139
2  2013-12-12  24.616
3  2013-12-13  25.369
4  2013-12-16  25.739
5  2013-12-17  25.245
6  2013-12-18  25.369
7  2013-12-19  25.265
8  2013-12-20  25.466
9  2013-12-23  25.324
10 2013-12-24  25.389
11 2013-12-26  25.275
12 2013-12-27  24.123
13 2013-12-30  23.967
14 2013-12-31  24.422
```

```
[17]: # Transformation of the series into a dataframe
series2 = df2['Close'].to_numpy()
X, Y = transformSeriesToDataset(series2, NumberOfElements = 10)
```

```
[18]: # Prediction
y_predict = lin_model.predict(X)
```

```
[20]: # Calculation of the mean squared error
MSE = mean_squared_error(Y, y_predict)
print("Test data evaluation: MSE =" + str(MSE))
```

Test data evaluation: MSE =0.8404990422436212

```
[22]: # The comparison between the real value and the prediction
df2_predictions = pd.DataFrame({'Real_Value':Y, 'Prediction':y_predict, 'Diff':
    ↪ Y-y_predict})
df2_predictions = df2_predictions.reset_index(drop = True)
df2_predictions.tail()
```

```
[22]:
```

	Real_Value	Prediction	Diff
973	47.346	47.120376	0.225624
974	47.406	47.283819	0.122181
975	46.358	47.519122	-1.161122
976	46.269	46.375744	-0.106744
977	45.670	46.360802	-0.690802

```
[58]: lin_model.predict(X[-2:])
```

```
[58]: array([46.37574405, 46.36080185])
```

```
[65]: # Generates a df to predict the next one
last_last_array=np.array(df2_predictions["Prediction"][-11:-1])
last_array=np.array(df2_predictions["Prediction"][-10:])
last_df = last_last_array,last_array
```

```
[68]: last_df
```

```
[68]: (array([50.83182221, 48.34180402, 47.51436499, 47.33626258, 46.71770821,
    47.94578569, 47.12037621, 47.28381897, 47.51912232, 46.37574405]),
    array([48.34180402, 47.51436499, 47.33626258, 46.71770821, 47.94578569,
    47.12037621, 47.28381897, 47.51912232, 46.37574405, 46.36080185]))
```

```
[74]: # Next one prediction
next_value = lin_model.predict(last_df)
print("The next value is:",round(next_value[1],2))
```


The next value is: 46.44

5 Step 4: Model Exploration / Explanation

Explain the model's predictions to someone who is unfamiliar with machine learning. Use charts/tables/etc to explain why they should (or perhaps should not?) rely on these predictions.

```
[23]: plt.figure(figsize=(20,10))
sns.set_style("whitegrid")
ax = sns.lineplot(data = df2_predictions["Real_Value"], label = "Real Value")
sns.lineplot(data = df2_predictions["Prediction"], label = "Prediction")
ax.set_title("Stock Price - Real Value VS Prediction Value", fontsize=25)
ax.set_ylabel("Price ($)", fontsize=25)
```

```
[23]: Text(0, 0.5, 'Price ($)')
```



This is the comparison of the real values vs the predicted values for the American Airlines stock “Close” price